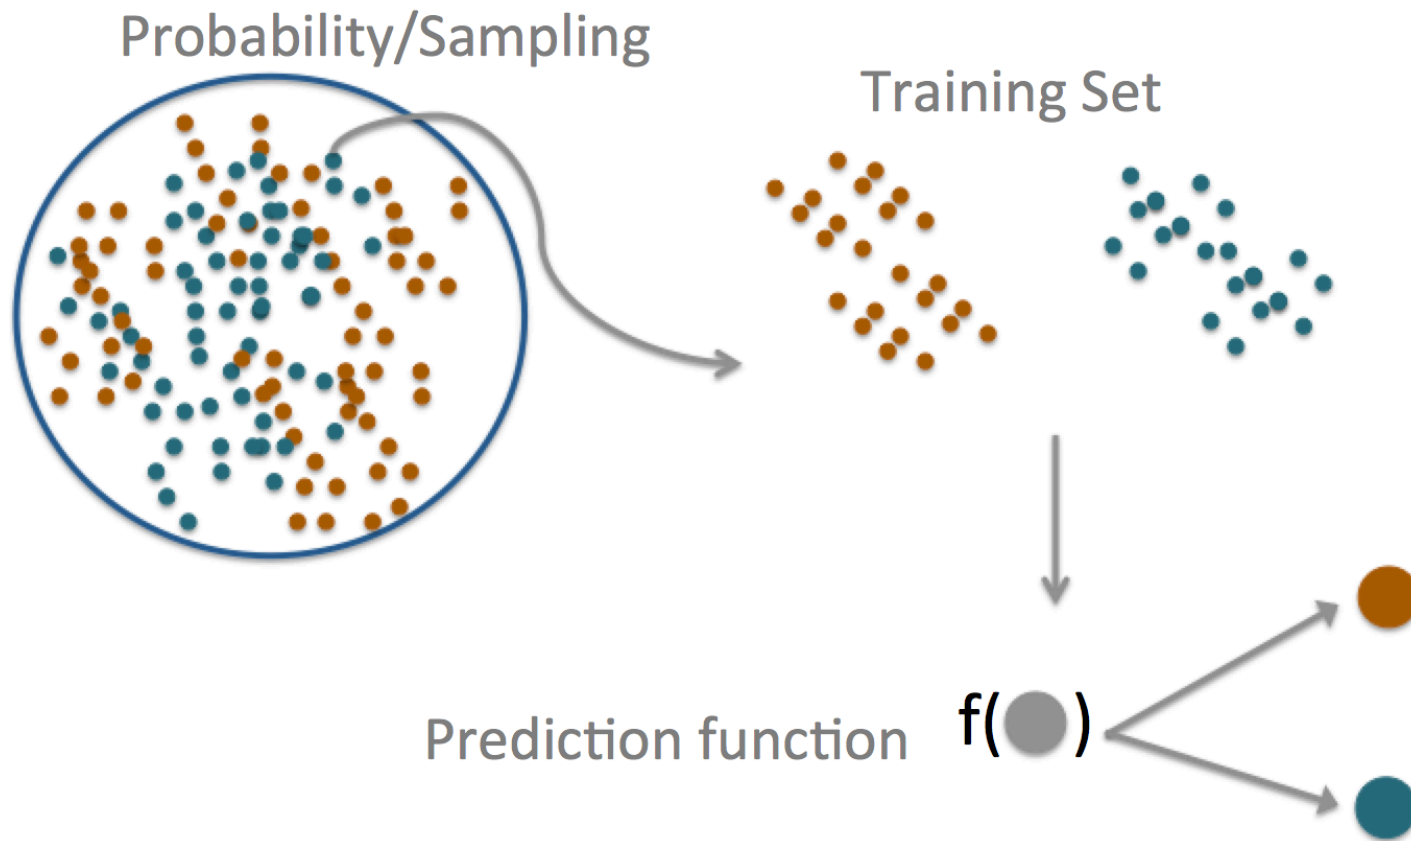




# What is prediction?

Jeffrey Leek  
Johns Hopkins Bloomberg School of Public Health

# The central dogma of prediction




# What can go wrong

## BIG DATA

### The Parable of Google Flu: Traps in Big Data Analysis

David Lazer,<sup>1,2\*</sup> Ryan Kennedy,<sup>1,3,4</sup> Gary King,<sup>3</sup> Alessandro Vespignani<sup>5,6,3</sup>

In February 2013, Google Flu Trends (GFT) made headlines but not for a reason that Google executives or the creators of the flu tracking system would have hoped. *Nature* reported that GFT was predicting more than double the proportion of doctor visits for influenza-like illness (ILI) than the Centers for Disease Control and Preven-



Large errors in flu prediction were largely avoidable, which offers lessons for the use of big data.

run ever since, with a few changes announced in October 2013 (10, 15).

Although not widely reported until 2013, the new GFT has been persistently overestimating flu prevalence for a much longer time. GFT also missed by a very large margin in the 2011–2012 flu season and has missed high for 100 out

<http://www.sciencemag.org/content/343/6176/1203.full.pdf>

# Components of a predictor

question -> input data -> features -> algorithm -> parameters -> evaluation

# SPAM Example

question -> input data -> features -> algorithm -> parameters -> evaluation

## **Start with a general question**

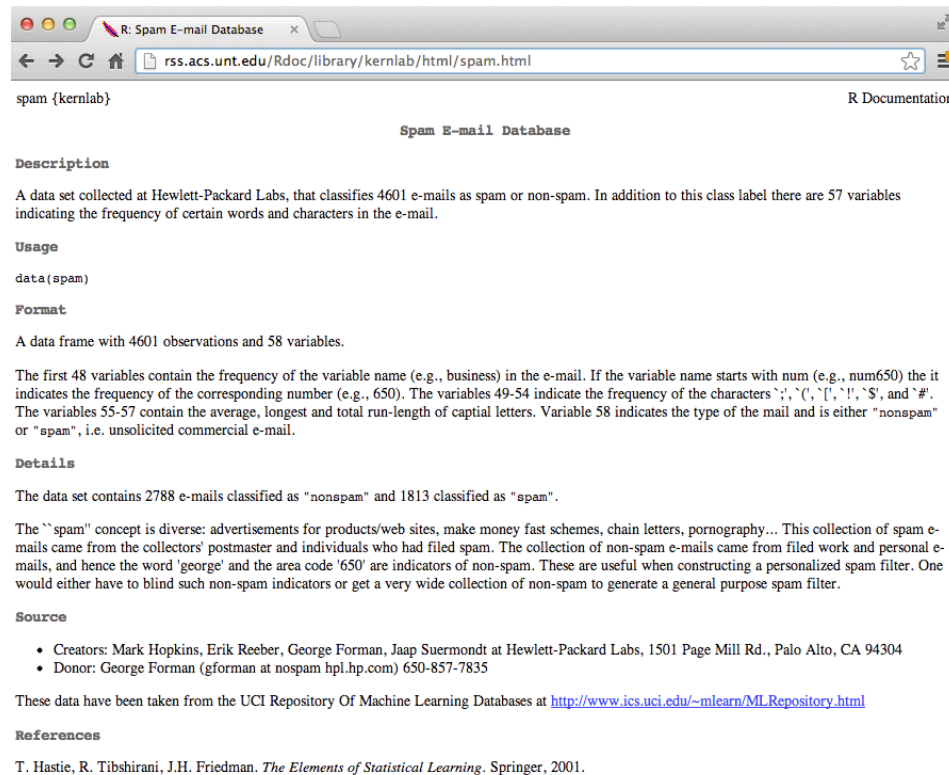
Can I automatically detect emails that are SPAM that are not?

## **Make it concrete**

Can I use quantitative characteristics of the emails to classify them as SPAM/HAM?

# SPAM Example

question -> **input data** -> features -> algorithm -> parameters -> evaluation



The screenshot shows a web browser window with the title "R: Spam E-mail Database". The address bar displays the URL [rss.acs.unt.edu/Rdoc/library/kernlab/html/spam.html](http://rss.acs.unt.edu/Rdoc/library/kernlab/html/spam.html). The page content is titled "Spam E-mail Database" and includes sections for Description, Usage, Format, Details, Source, and References.

**Description**

A data set collected at Hewlett-Packard Labs, that classifies 4601 e-mails as spam or non-spam. In addition to this class label there are 57 variables indicating the frequency of certain words and characters in the e-mail.

**Usage**

```
data(spam)
```

**Format**

A data frame with 4601 observations and 58 variables.

The first 48 variables contain the frequency of the variable name (e.g., business) in the e-mail. If the variable name starts with num (e.g., num650) the it indicates the frequency of the corresponding number (e.g., 650). The variables 49-54 indicate the frequency of the characters `', `(', `[', `!', `\$', and `#'. The variables 55-57 contain the average, longest and total run-length of capital letters. Variable 58 indicates the type of the mail and is either "nonspam" or "spam", i.e. unsolicited commercial e-mail.

**Details**

The data set contains 2788 e-mails classified as "nonspam" and 1813 classified as "spam".

The "spam" concept is diverse: advertisements for products/web sites, make money fast schemes, chain letters, pornography... This collection of spam e-mails came from the collectors' postmaster and individuals who had filed spam. The collection of non-spam e-mails came from filed work and personal e-mails, and hence the word 'george' and the area code '650' are indicators of non-spam. These are useful when constructing a personalized spam filter. One would either have to blind such non-spam indicators or get a very wide collection of non-spam to generate a general purpose spam filter.

**Source**

- Creators: Mark Hopkins, Erik Reeber, George Forman, Jaap Suermondt at Hewlett-Packard Labs, 1501 Page Mill Rd., Palo Alto, CA 94304
- Donor: George Forman (gforman at nospam hpl.hp.com) 650-857-7835

These data have been taken from the UCI Repository Of Machine Learning Databases at <http://www.ics.uci.edu/~mlern/MLRepository.html>

**References**

T. Hastie, R. Tibshirani, J.H. Friedman. *The Elements of Statistical Learning*. Springer, 2001.

<http://rss.acs.unt.edu/Rdoc/library/kernlab/html/spam.html>

# SPAM Example

question -> input data -> features -> algorithm -> parameters -> evaluation

Dear Jeff,

Can you send me your address so I can send you the invitation?

Thanks,

Ben

# SPAM Example

question -> input data -> features -> algorithm -> parameters -> evaluation

Dear Jeff,

Can **you** send me your address so I can send **you** the invitation?

Thanks,

Ben

Frequency of you =  $2/17 = 0.118$



# SPAM Example

question -> input data -> **features** -> algorithm -> parameters -> evaluation

```
library(kernlab)
data(spam)
head(spam)
```

	make	address	all	num3d	our	over	remove	internet	order	mail	receive	will	people	report	addresses
1	0.00	0.64	0.64	0	0.32	0.00	0.00	0.00	0.00	0.00	0.00	0.64	0.00	0.00	0.00
2	0.21	0.28	0.50	0	0.14	0.28	0.21	0.07	0.00	0.94	0.21	0.79	0.65	0.21	0.14
3	0.06	0.00	0.71	0	1.23	0.19	0.19	0.12	0.64	0.25	0.38	0.45	0.12	0.00	1.75
4	0.00	0.00	0.00	0	0.63	0.00	0.31	0.63	0.31	0.63	0.31	0.31	0.31	0.00	0.00
5	0.00	0.00	0.00	0	0.63	0.00	0.31	0.63	0.31	0.63	0.31	0.31	0.31	0.00	0.00
6	0.00	0.00	0.00	0	1.85	0.00	0.00	1.85	0.00	0.00	0.00	0.00	0.00	0.00	0.00


  

	free	business	email	you	credit	your	font	num000	money	hp	hpl	george	num650	lab	labs	telnet
1	0.32	0.00	1.29	1.93	0.00	0.96	0	0.00	0.00	0	0	0	0	0	0	0
2	0.14	0.07	0.28	3.47	0.00	1.59	0	0.43	0.43	0	0	0	0	0	0	0
3	0.06	0.06	1.03	1.36	0.32	0.51	0	1.16	0.06	0	0	0	0	0	0	0
4	0.31	0.00	0.00	3.18	0.00	0.31	0	0.00	0.00	0	0	0	0	0	0	0

# SPAM Example

question -> input data -> features -> **algorithm** -> parameters -> evaluation

```
plot(density(spam$your[spam$type=="nonspam"]),  
     col="blue",main="",xlab="Frequency of 'your'")  
lines(density(spam$your[spam$type=="spam"]),col="red")
```



# SPAM Example

question -> input data -> features -> **algorithm** -> parameters -> evaluation


## Our algorithm

- Find a value  $C$ .
- **frequency of 'your' > C** predict "spam"

# SPAM Example

question -> input data -> features -> algorithm -> **parameters** -> evaluation

```
plot(density(spam$your[spam$type=="nonspam"]),  
     col="blue",main="",xlab="Frequency of 'your'")  
lines(density(spam$your[spam$type=="spam"]),col="red")  
abline(v=0.5,col="black")
```



# SPAM Example

question -> input data -> features -> algorithm -> parameters -> **evaluation**

```
prediction <- ifelse(spam$your > 0.5, "spam", "nonspam")  
table(prediction, spam$type) / length(spam$type)
```

prediction	nonspam	spam
nonspam	0.4590	0.1017
spam	0.1469	0.2923

Accuracy  $\approx 0.459 + 0.292 = 0.751$