



Plotting predictors

Jeffrey Leek
Johns Hopkins Bloomberg School of Public Health

Example: predicting wages



Image Credit <http://www.caahs-media.org/the-high-cost-of-low-wages>

Data from: [ISLR package](#) from the book: [Introduction to statistical learning](#)

Example: Wage data

```
library(ISLR); library(ggplot2); library(caret);  
data(Wage)  
summary(Wage)
```

year	age	sex	maritl	race
Min. :2003	Min. :18.0	1. Male :3000	1. Never Married: 648	1. White:2480
1st Qu.:2004	1st Qu.:33.8	2. Female: 0	2. Married :2074	2. Black: 293
Median :2006	Median :42.0		3. Widowed : 19	3. Asian: 190
Mean :2006	Mean :42.4		4. Divorced : 204	4. Other: 37
3rd Qu.:2008	3rd Qu.:51.0		5. Separated : 55	
Max. :2009	Max. :80.0			

education	region	jobclass	health
1. < HS Grad :268	2. Middle Atlantic :3000	1. Industrial :1544	1. <=Good : 858
2. HS Grad :971	1. New England : 0	2. Information:1456	2. >=Very Good:2142
3. Some College :650	3. East North Central: 0		
4. College Grad :685	4. West North Central: 0		
5. Advanced Degree:426	5. South Atlantic : 0		
	6. East South Central: 0		

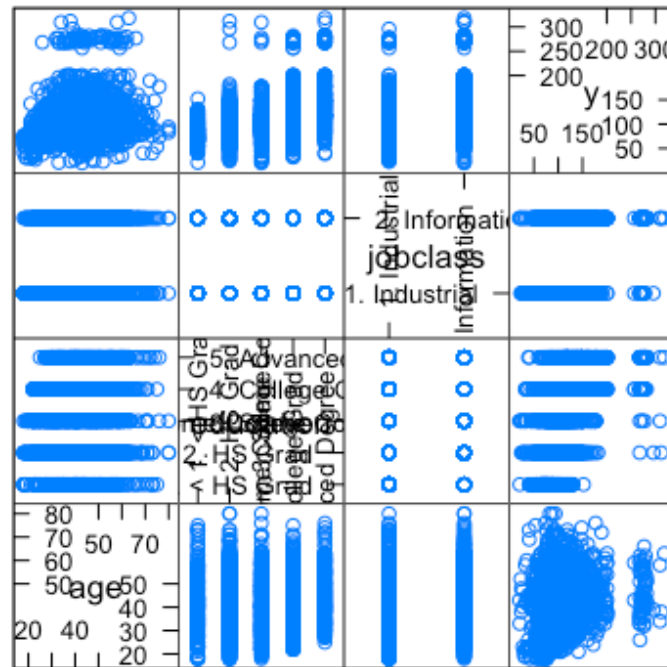
Get training/test sets

```
inTrain <- createDataPartition(y=Wage$wage,  
                                p=0.7, list=FALSE)  
training <- Wage[inTrain,]  
testing  <- Wage[-inTrain,]  
dim(training); dim(testing)
```

```
[1] 898 12
```

Feature plot (*caret* package)


```
featurePlot(x=training[,c("age", "education", "jobclass")],  
            y = training$wage,  
            plot="pairs")
```



Scatter Plot Matrix


Qplot (*ggplot2* package)

```
qplot(age,wage,data=training)
```




Qplot with color (*ggplot2* package)

```
qplot(age,wage,colour=jobclass,data=training)
```



Add regression smoothers (*ggplot2* package)

```
qq <- qplot(age, wage, colour=education, data=training)
qq + geom_smooth(method='lm', formula=y~x)
```




cut2, making factors (*Hmisc* package)

```
cutWage <- cut2(training$wage,g=3)
table(cutWage)
```

```
cutWage
[ 20.1, 91.7) [ 91.7,118.9) [118.9,318.3]
           704           725           673
```



Boxplots with cut2

```
p1 <- ggplot(cutWage, age, data=training, fill=cutWage,  
             geom=c("boxplot"))  
p1
```



Boxplots with points overlaid

```
p2 <- ggplot(cutWage, age, data=training, fill=cutWage,  
             geom=c("boxplot", "jitter"))  
grid.arrange(p1, p2, ncol=2)
```



Tables

```
t1 <- table(cutWage,training$jobclass)
t1
```


```
cutWage      1. Industrial 2. Information
[ 20.1, 91.7)          437          267
[ 91.7,118.9)          365          360
[118.9,318.3]          263          410
```

```
prop.table(t1,1)
```

```
cutWage      1. Industrial 2. Information
[ 20.1, 91.7)          0.6207          0.3793
[ 91.7,118.9)          0.5034          0.4966
[118.9,318.3]          0.3908          0.6092
```

Density plots

```
qplot(wage, colour=education, data=training, geom="density")
```



Notes and further reading

- Make your plots only in the training set
 - Don't use the test set for exploration!
- Things you should be looking for
 - Imbalance in outcomes/predictors
 - Outliers
 - Groups of points not explained by a predictor
 - Skewed variables
- [ggplot2 tutorial](#)
- [caret visualizations](#)