



# Unsupervised prediction

Jeffrey Leek, Assistant Professor of Biostatistics  
Johns Hopkins Bloomberg School of Public Health

# Key ideas

- Sometimes you don't know the labels for prediction
- To build a predictor
  - Create clusters
  - Name clusters
  - Build predictor for clusters
- In a new data set
  - Predict clusters

# Iris example ignoring species labels


```
data(iris); library(ggplot2)
inTrain <- createDataPartition(y=iris$Species,
                                p=0.7, list=FALSE)

training <- iris[inTrain,]
testing  <- iris[-inTrain,]
dim(training); dim(testing)
```

```
[1] 45  5
```

# Cluster with k-means

```
kMeans1 <- kmeans(subset(training,select=-c(Species)),centers=3)
training$clusters <- as.factor(kMeans1$cluster)
qplot(Petal.Width,Petal.Length,colour=clusters,data=training)
```



# Compare to real labels

```
table(kMeans1$cluster, training$Species)
```

	setosa	versicolor	virginica
1	0	1	23
2	35	0	0
3	0	34	12

# Build predictor

```
modFit <- train(clusters ~.,data=subset(training,select=-c(Species)),method="rpart")  
table(predict(modFit,training),training$Species)
```

	setosa	versicolor	virginica
1	0	0	21
2	35	0	0
3	0	35	14

# Apply on test

```
testClusterPred <- predict(modFit,testing)
table(testClusterPred ,testing$Species)
```

```
testClusterPred setosa versicolor virginica
1          0          0          13
2         15          0          0
3          0         15          2
```

# Notes and further reading

- The `cl_predict` function in the `clue` package provides similar functionality
- Beware over-interpretation of clusters!
- This is one basic approach to [recommendation engines](#)
- [Elements of statistical learning](#)
- [Introduction to statistical learning](#)